

Great. I'll prepare a structured report covering your questions and answers about the Hessian matrix in the 3-parameter Lie algebra space, including its derivation, properties, and implications for optimization quality.

I'll organize the report into clearly labeled sections, explaining:

- How the Hessian is derived from the residual
- How it guides point selection
- How it evaluates optimality
- Its relationship to the information matrix
- The role of left vs. right perturbations
- Why the Gauss–Newton approximation is justified
- How the condition number influences solution quality

I'll let you know when the report is ready for review.

## Hessian in 3D Rigid-Body Optimization Using Lie Algebra ( $\omega$ ) Representation

**Introduction:** In 3D rigid-body least-squares problems (e.g. point cloud registration or pose estimation), analyzing the Hessian of the cost function provides deep insights into convergence and solution quality. We consider rotation parameterized by a 3-parameter Lie algebra vector  $\omega$  ( $\mathfrak{so}(3)$  axis-angle increments) and possibly translation. For a given cost function  $E(R, t) = \sum_i \|y_i - R x_i - t\|^2$ , we focus on the **rotation part**. We derive the Gauss–Newton Hessian  $H = J^T J$  with respect to  $\omega$ , discuss how the point configuration affects conditioning, and relate Hessian properties to optimality and uncertainty. We also compare **left vs. right perturbation** conventions and justify the Gauss–Newton approximation. Throughout, we assume small rotation increments (exp-map linearization) and treat the residuals  $r_i = y_i - R x_i - t$ . The target audience is advanced students or practitioners in vision/robotics/SLAM, so we emphasize clear math notation and intuitive explanations.

### 1. Derivation of the Gauss–Newton Hessian $J^T J$ for Rotation

**Jacobian w.r.t.  $\omega$ :** Using a left perturbation model (update  $R = \exp([\omega]_\times) R_0$ ), the residual for point  $i$  is  $r_i = y_i - R x_i - t$ . A small rotation  $\delta\omega$  induces  $R x_i \approx R_0 x_i + [\delta\omega]_\times (R_0 x_i)$ . Noting  $a \times b = -b \times a$ , we have  $[\delta\omega]_\times (R_0 x_i) = -[R_0 x_i]_\times \delta\omega$ . Thus the **Jacobian** of the residual w.r.t. the Lie vector  $\omega$  is:

$$J_i = \left. \frac{\partial r_i}{\partial \omega} \right|_{\omega=0} = -[R_0 x_i]_\times,$$

a  $3 \times 3$  matrix. Geometrically,  $[R_0 x_i]_\times$  is the skew-symmetric cross-product matrix of the vector  $R_0 x_i$ . This matches the standard result that the angular

error derivative is the negative cross product with the current rotated point. (If a **right perturbation**  $R = R_0 \exp([\omega]_{\times})$  were used instead, the Jacobian would appear as  $-[R_0 x_i]_{\times} R_0$ , complicating the form by involving  $R_0$ . Left perturbation is preferred since it yields a simpler, frame-invariant Jacobian and is used in most SLAM/ICP libraries for consistency.)

**Gauss–Newton Hessian:** The approximate Hessian for rotation (Gauss–Newton) is  $H = \sum_i J_i^T J_i = \sum_i ([R_0 x_i]_{\times})^T [R_0 x_i]_{\times}$ . Using the vector identity  $[v]_{\times}^T [v]_{\times} = \|v\|^2 I - v v^T$ , this becomes:

$$H = \sum_i \left( \|R_0 x_i\|^2 I_{3 \times 3} - (R_0 x_i)(R_0 x_i)^T \right).$$

Since rotation preserves lengths,  $\|R_0 x_i\| = \|x_i\|$ . If we assume the current estimate  $R_0$  is identity (or we analyze at the optimum where residuals are minimized), we can drop  $R_0$  for intuition. Thus an equivalent form is:

$$H = \sum_i \left( \|x_i\|^2 I - x_i x_i^T \right). \quad (1)$$

This matrix has a clear physical interpretation: it is essentially the **inertia matrix** of the point set  $\{x_i\}$  about the origin. Each term  $\|x_i\|^2 I - x_i x_i^T$  resembles a point mass’s contribution to a rotation inertia tensor. Expanding for insight: if  $x_i = [x, y, z]^T$ , the term adds  $\|x_i\|^2$  to each diagonal and subtracts the outer product  $x_i x_i^T$ . Summing over points yields a symmetric  $3 \times 3$  matrix **H** that is positive semi-definite (as  $J^T J$  always is). We will analyze its rank, eigenvalues, and definiteness next.

*Derivation check:* If only one point  $x$  is used,  $H = \|x\|^2 I - x x^T$ . This has one zero eigenvalue (along the direction of  $x$ ) and two positive eigenvalues equal to  $\|x\|^2$ . Intuitively, a single point constrains rotation only orthogonal to the line of that point (rotating about the point’s axis doesn’t change error). With many points, if they are in “general position,” **H** will be full rank; otherwise it will have small or zero eigenvalues for unconstrained rotation axes, as we discuss below.

## 2. Point Selection and Hessian Conditioning

Equation (1) reveals that **H** depends purely on the spatial distribution of the points  $x_i$ . This guides how to select or arrange points to ensure good conditioning (i.e. well-behaved, non-singular Hessian):

- **Well-conditioned (optimal) configurations:** To fully constrain all 3 rotation axes, the points should not lie in any lower-dimensional subspace. An ideal configuration is spread out in 3D such that no axis passes through all points without moving them. For example, four non-coplanar points (like vertices of a tetrahedron) or a spherical distribution around the origin

yields a strong, isotropic Hessian. In such cases,  $\mathbf{H}$  is positive definite with relatively balanced eigenvalues (no eigenvalue is too small compared to others). The condition number (ratio  $\lambda_{\max}/\lambda_{\min}$ ) is near 1 if points are symmetrically distributed. Intuitively, rotations about any axis will significantly move some points, producing measurable error changes.

- **Degenerate configurations:** If points lie in a plane or a line (or very close to it),  $\mathbf{H}$  becomes ill-conditioned or rank-deficient. For instance, if all  $x_i$  lie exactly on a line through the origin, rotation about that line produces no change in  $Rx_i$ , so the Hessian has an eigenvalue of 0 in that direction. Indeed, our formula gives  $Hd = 0$  if  $d$  is along the line (because  $x_i^T d = \|x_i\| \|d\| \cos \theta = \|x_i\| \|d\|$  for points on the line, making each term  $\|x_i\|^2 d - (x_i^T d)x_i = 0$ ). In this **collinear** case,  $H$  has rank 2 and is singular. Likewise, if points are coplanar through the origin, rotation about the normal to that plane is fully constrained (points move within the plane), but rotations about axes in the plane still move points albeit all in the plane. In fact, if points span a plane,  $H$  will typically be full rank (all rotations affect them) but one eigenvalue will be smaller – the axis normal to the plane tends to be a **principal axis** with a large moment (easy to rotate about), and the weakest-constrained rotation is about an axis in the plane (if the point distribution in-plane is symmetric, two eigenvalues correspond to in-plane principal moments, one could be smaller).
- **Examples:** Consider three non-collinear points on a plane through the origin. They span the plane, so no rotation leaves all points fixed except the trivial identity;  $H$  is positive-definite. However, one eigenvalue will correspond to rotating around the plane’s normal (which moves points in the plane – usually a larger inertia), and the smallest eigen corresponds to rotating about some in-plane axis (which still moves points out of the plane). By contrast, if one point set is extremely near a line or plane, the Hessian will be ill-conditioned with one eigenvalue others, indicating a “nearly degenerate” case.

**Guiding point selection:** The goal is to maximize the “spread” of points in all directions so that  $H$  is well-conditioned. In practice, one can measure the **condition number** or the smallest eigenvalue of  $H$  as a criterion. A very small  $\lambda_{\min}$  (or large condition number) indicates degeneracy or weak observability of some rotation axis. An optimal set of points would maximize  $\lambda_{\min}$  (making Hessian closer to isotropic). For instance, placing points at large distances along orthogonal directions (x, y, z axes) will increase the diagonal terms and ensure off-diagonal terms cancel out less, yielding a strong Hessian. On the other hand, points clustered along one direction (even if separated) lead to a near-singular  $H$ . Some algorithms actively **resample or choose correspondences** to improve conditioning, aiming for a Hessian with condition number close to 1.

In summary, Equation (1) highlights that **diverse point geometries** (not all points on the same line or plane) are necessary for a well-conditioned Hessian. Degenerate configurations reduce the rank of  $H$ , causing flat cost directions (the

optimization will have a continuum of solutions or very slow curvature in some directions). By ensuring points span 3D space, we avoid such issues.

### 3. What the Hessian Says About the Optimum’s Quality

At the optimum (or at a candidate solution), the Hessian matrix  $H$  reveals the **curvature** of the cost function in different rotational directions:

- **Positive definiteness:** A Hessian  $H$  that is positive definite (all eigenvalues  $> 0$ ) indicates a local minimum in all directions. In a least-squares pose problem without degeneracy,  $H$  should be PD at the true optimum (meaning the solution is unique and stable for small perturbations). If  $H$  has a zero eigenvalue, it means the error function is flat in some direction – either a continuum of solutions or an unobservable rotation axis. For example, if our points were collinear, any rotation about that line yields the same error (a whole set of optima). In such cases, the optimum is not well-defined (the cost is minimized along a “valley”). Therefore, checking that  $H$  is non-singular (or has all large positive eigenvalues) is crucial for a well-defined optimum. A negative eigenvalue of the true Hessian (not the Gauss–Newton approximation) would indicate a saddle or maximum, but in Gauss–Newton  $J^T J$  cannot be negative – it’s at least semi-definite. So, in practice for least squares, non-PD usually means singular or nearly so.
- **Eigenvalues and curvature:** The magnitude of an eigenvalue reflects how sharply the error increases for a small rotation in the corresponding eigenvector direction. A **large eigenvalue** means a steep curvature (the cost grows quickly in that direction, so the solution is stiffly constrained for that rotation axis). A **small eigenvalue** means a nearly flat direction (the cost grows slowly, indicating that the optimum is “sloppy” or ill-constrained in that direction). The ratio between the largest and smallest eigenvalues (condition number) quantifies the anisotropy of the curvature. A condition number near 1 means the error surface is nearly spherical (equally curved in all directions) – the best scenario for stable convergence. A very large condition number indicates an elongated error bowl, implying the solution is very sensitive to noise or initial errors in the weak direction.
- **Determinant and volume:** The determinant of  $H$  is the product of its eigenvalues ( $\lambda_1 \lambda_2 \lambda_3$ ). In estimation theory this is related to the volume of the uncertainty ellipsoid. A larger determinant (all else equal) means a smaller uncertainty volume. Maximizing  $\det(H)$  (equivalently minimizing the volume of the confidence ellipsoid) is sometimes a criterion for experiment design. In our context, if one eigenvalue is zero,  $\det(H) = 0$  indicating infinite volume (complete uncertainty in one rotation).
- **Confidence and covariance:** If the measurement noise is small, one can approximate the **covariance** of the estimate as  $\Sigma \approx \sigma^2 H^{-1}$  (where  $\sigma^2$  is the noise variance per residual). Thus, eigenvalues of  $H$  inversely relate to the variance of the estimated rotation around certain axes. A small

eigenvalue (poor curvature) corresponds to high variance (uncertainty) in that rotation angle, meaning the optimum could shift significantly with slight noise changes. Conversely, large eigenvalues correspond to low variance, i.e., that aspect of the rotation is well-pinned by the data.

In summary, a “good” optimum is indicated by a well-conditioned, positive-definite Hessian: all eigenvalues positive and not wildly disparate. This ensures the solution is unique (no flat directions), and that the error surface is steep enough in all directions for the numerical optimizer to find it and for the estimate to be robust. Practically, one might inspect  $H$  at the found solution to confirm it is PD and maybe check the smallest eigenvalue or condition number as a **quality metric**. If  $H$  is near-singular, it signals that the solution (or model) has an ambiguity – additional constraints or points may be needed to improve confidence.

#### 4. Gauss–Newton Hessian vs. Fisher Information Matrix

In least-squares problems under Gaussian noise assumptions, the Gauss–Newton approximation  $H = J^T J$  is closely tied to the **Fisher Information Matrix (FIM)**. In fact, for a correctly-modeled system,  $J^T J$  (scaled by the noise precision  $1/\sigma^2$ ) is the Fisher information matrix of the parameter estimate. The Fisher information  $I(\theta)$  is defined as the expected curvature of the log-likelihood at the optimum:

- In our case, each residual  $r_i = y_i - Rx_i - t$  is assumed to have noise (say iid Gaussian with variance  $\sigma^2$ ). The log-likelihood of the data given parameters is  $-\frac{1}{2\sigma^2} \sum_i \|r_i\|^2 + \text{const}$ . The Hessian of the negative log-likelihood w.r.t. the parameters is  $\frac{1}{\sigma^2} \sum_i (\nabla r_i^T \nabla r_i + r_i \nabla^2 r_i)$ . The Gauss–Newton method drops the second term (which is zero at the optimum if residuals  $\rightarrow 0$ ). So at the optimum (or under small residual assumption),  $\nabla^2(-\log L) \approx \frac{1}{\sigma^2} J^T J$ . The **Fisher information matrix** is defined as  $I = E[\nabla^2(-\log L)]$ , which under those assumptions equals  $\frac{1}{\sigma^2} J^T J$  evaluated at the true parameter. Thus  $J^T J$  is essentially the observed information matrix.
- **Equivalence:** It has been shown that for models in the exponential family (Gaussians for least squares, etc.), the Gauss–Newton Hessian is equivalent to the Fisher information. In simpler terms,  $J^T J$  provides an estimate of how much “information” the data carry about the rotation parameters. Large entries/eigenvalues in  $J^T J$  mean the data are informative (e.g., many points at large radii constrain rotation strongly), whereas a singular  $J^T J$  means some parameters are unobservable from the data.
- **Optimality and Uncertainty:** Because of this equivalence, assessing optimality and uncertainty via  $H = J^T J$  is meaningful. In estimation theory, the Cramér–Rao bound tells us the covariance lower bound is  $I^{-1}$ . So if we compute  $H^{-1}$  (when possible), we get an approximation of the covariance of the rotation estimate. For instance, if one axis has significantly larger variance, it will show up as a small eigenvalue in  $H$ . This links

back to Section 3’s discussion – a well-conditioned Hessian (information matrix) implies a **low-uncertainty estimate** in all directions, whereas a poorly conditioned one implies the estimate has a large uncertainty in some direction.

- **Interpreting values:** Each diagonal element of  $H$  roughly corresponds to the total squared leverage of all points on one axis of rotation, and off-diagonals indicate correlation between axes. In practice, least-squares SLAM and bundle adjustment frameworks use the information matrix for smoothing and marginalization. For example, in SLAM graph optimization, the accumulated Hessian (normal matrix) of the system represents the information matrix of the map and robot pose estimates. Examining it can tell which directions in state space are constrained by measurements.

In summary, **Gauss–Newton’s Hessian is not just a numerical artifact; it’s the empirical information matrix of the problem.** This justifies using  $H$  to gauge solution quality: a high information (well-conditioned  $H$ ) means the data strongly pin down the rotation (low uncertainty), while a low-rank or ill-conditioned  $H$  means the data lack information in some rotational degree (leading to ambiguity or high uncertainty). This connection also explains why second-order optimization methods (like Gauss–Newton) and statistical estimation theory converge: both are effectively analyzing the curvature of the same likelihood surface.

## 5. Left vs. Right Perturbation in Lie Algebra Rotation Optimization

When optimizing over rotations with Lie algebra parameterizations, one must choose how to apply small updates ( $\delta\omega$ ) to the current rotation estimate  $R$ . The two common conventions are:

- **Left perturbation:**  $R_{\text{new}} = \exp([\delta\omega]_{\times}) R_{\text{old}}$ . Here the increment is applied on the left, in the **global frame**. Intuitively, we “rotate the current rotation” by a small angle in the fixed world coordinate system.
- **Right perturbation:**  $R_{\text{new}} = R_{\text{old}} \exp([\delta\omega]_{\times})$ . The increment is applied on the right, i.e. in the **local frame** of the moving body. This rotates the coordinate frame attached to the body by a small angle.

While both are valid (and ultimately represent the same group tangent update), **most implementations use left perturbation** for rotation optimization. There are several reasons for this preference:

(a) **Simpler Jacobian and residual form:** As derived above, left perturbation yields  $J = -[Rx]_{\times}$  for the point residual  $y - Rx$ . This form is nice because it doesn’t explicitly depend on the current rotation  $R$ ; it naturally lives in the world frame. In contrast, right perturbation produces a Jacobian of form  $J = -R[x]_{\times}$  (when working out the derivative, an extra  $R$  appears because the local-axis perturbation must be mapped to world changes). This can be inconvenient: it

mixes the current state into the Jacobian, and one must be careful to transform the parameter update into the correct frame. Left perturbation effectively treats the residual  $r = y - Rx$  as living in the world frame, so a small rotation  $\delta\omega$  (in world axes) changes  $r$  by  $-[Rx]_{\times}\delta\omega$ . This **residual linearization is simpler**. Many SLAM/ICP libraries adopt left-increment conventions so that the math is cleaner and less error-prone.

**(b) Consistency with group composition:** In problems with multiple poses or composing transformations, left perturbations correspond to the standard group action on state. For example, if a pose is represented by a transformation matrix  $T = [R|t]$ , a left perturbation  $T \leftarrow \exp(\delta\xi)T$  (with  $\delta\xi = [\delta\omega, \delta t]$  in  $\text{se}(3)$ ) keeps the mathematical form of compounding consistent. It essentially says we correct the pose by a small motion in the world frame. This is coherent with how measurement functions are often written (world points projected via the pose). Right perturbation would mean adjusting the pose in its local frame, which can be counter-intuitive when mixing with global measurements.

**(c) Invariance properties:** Left-invariant error formulations (common in invariant filtering and optimization) result from left perturbations. They ensure that if you express errors in the global frame, the derivatives do not depend on the current estimate (which is a desirable property in Gaussian-Newton and EKF linearizations — the Jacobian becomes constant for certain models). Right perturbation would lead to right-invariant errors (errors fixed in body frame), which in some contexts (like IMU preintegration) are used, but for point registration, left-invariant tends to align with how the measurements are defined (e.g. world coordinates of points). In practical terms, applying left updates often avoids the need for extra adjoint transformations in the Jacobian.

**(d) Alignment with existing libraries:** Popular libraries such as GTSAM, Ceres (with its local parameterization for quaternions), Sophus etc., typically implement the **manifold update for  $\text{SO}(3)$**  as left multiplication by  $\exp(\delta\omega)$ . This means the incremental solving ( $R \leftarrow \exp(\delta\omega)R$ ) is done with left perturbation. Adhering to this convention makes it easier to integrate with these tools and with standard equations in literature. It also ensures consistency when combining rotation with translation updates in  $\text{SE}(3)$ . For example, if the full pose is  $T = (R, t)$ , one often writes  $T_{\text{new}} = \exp([\delta\omega, \delta t]_{\wedge})T_{\text{old}}$ , which implies left multiplication on both rotation and translation — a common choice in optimization frameworks.

In summary, left perturbation is preferred because it yields a simpler linearization of residuals (global frame derivatives), making the Gauss–Newton (or Levenberg–Marquardt) steps easier to derive and implement correctly. It also meshes well with the prevailing formulations in SLAM/ICP where error is measured in a global coordinate system. Right perturbation can certainly be used (and in some robotic calibration problems it might be more natural), but one must then account for the rotation of the Jacobian into the correct frame. For the purposes of this report (deriving  $H = \sum_i (\|x_i\|^2 I - x_i x_i^T)$ ), we explicitly used left perturbation, as evidenced by the Jacobian  $-[Rx]_{\times}$  which came out cleanly.

## 6. Why Gauss–Newton Uses $J^\top J$ (Neglecting Second-Order Terms)

The Gauss–Newton method is an approximation to the true Newton’s method that is particularly suited for least-squares problems. It **ignores the second-order derivative of the residuals**, effectively using  $H \approx J^\top J$  instead of the full Hessian of the cost. The justification for this lies in the nature of residuals and the assumption of small errors:

- **Small residual assumption:** In least squares, the cost is  $F(\theta) = \frac{1}{2} \sum_i \|r_i(\theta)\|^2$ . Its gradient and Hessian are

$$\nabla F = \sum_i J_i^\top r_i, \quad H_{\text{full}} = \sum_i \left( J_i^\top J_i + \underbrace{\nabla^2 r_i r_i}_{\text{second-order term}} \right).$$

The term  $J_i^\top J_i$  is what Gauss–Newton keeps, and  $\nabla^2 r_i r_i$  is what it drops. If the residuals  $r_i$  are small (i.e., the current solution is close to the optimum and the model fits well), then the product of  $r_i$  with the second derivative of  $r_i$  is second-order small. Near the optimum of a well-behaved problem,  $r_i \approx 0$  for all  $i$  (exact zero if the model can perfectly fit the data). Thus, the neglected term is negligible in the vicinity of the solution. Gauss–Newton essentially linearizes the model  $r_i(\theta)$  and computes curvature of that linearized model’s error, assuming the nonlinearity’s curvature is minor when errors are already low.

- **Computational advantage:** Even when residuals are not extremely small, Gauss–Newton has a big practical advantage: one does **not need to compute second derivatives** of each residual function. Computing and assembling the full Hessian can be complex and costly, especially for large problems. In our rotation example, the residual  $r_i = y_i - Rx_i - t$  depends on  $R$  nonlinearly, but computing  $\nabla^2 r_i$  would involve second-order terms in  $\omega$  (which are related to the curvature of  $\text{SO}(3)$ ). Gauss–Newton bypasses this by only using  $J_i$ . This not only simplifies implementation but also guarantees that the approximate Hessian  $J^\top J$  is **positive semi-definite**. The true Hessian might not be PSD if the model has significant nonlinear effects (it could have negative eigenvalues indicating a locally non-convex direction). By dropping the troublesome second part, Gauss–Newton ensures we are descending in a convex quadratic approximation. The resulting  $H = J^\top J$  is PSD by construction (sum of outer products), avoiding the scenario of a negative curvature direction that could destabilize a Newton step.
- **Convergence behavior:** Neglecting second-order terms makes Gauss–Newton *less accurate* than full Newton per iteration, but often more robust. When the linear approximation of residuals is good (small errors, well-behaved Jacobian), Gauss–Newton converges quickly (usually quadratic convergence near the optimum). If residuals are not small,



Gauss–Newton can struggle (it might not decrease the cost if the Hessian approximation is too far off). In those cases, a **Levenberg–Marquardt** algorithm is commonly used, which interpolates between Gauss–Newton and gradient descent by adding a damping term to  $H$ . Essentially, LM handles cases where Gauss–Newton’s assumption doesn’t hold by stabilizing the Hessian.

- **When is Gauss–Newton valid?** Empirically, if you start with a decent initial guess (so residuals are moderate) and the problem is moderately nonlinear, Gauss–Newton tends to work well. Our rotation alignment problem is actually nicely behaved: the cost as a function of small rotation is fairly quadratic near the optimum (rotational misalignment causes a quadratic increase in error for small angles). Therefore,  $J^T J$  is a good approximation of the true Hessian near the optimum. If the rotation error is huge (say 180 degrees off), the residuals might be large and highly nonlinear in  $\omega$ , so Gauss–Newton could misbehave; one would then rely on damping or a better initialization (possibly using a global method or random rotations to find a basin of attraction).

**Effect on computation:** The Gauss–Newton Hessian  $J^T J$  being used means each iteration involves solving  $J^T J \delta\omega = J^T r$ . This is the normal equations system. This costs  $O(nm^2)$  to form and solve ( $n$  = number of points,  $m = 3$  for rotation), which is very efficient here. Computing the true Hessian would be similar complexity in this small case but for large-scale problems (like thousands of parameters) the difference is significant: computing second derivatives can be expensive and unnecessary. By using  $J^T J$ , we leverage the problem’s least-squares structure.

In summary, Gauss–Newton drops the second-order term because:

- **It’s small when we are near the minimum** (which is when we most care about accurate Hessian for convergence).
- **It ensures the Hessian approximation is PSD**, avoiding issues with indefiniteness.
- **It reduces computational and implementation complexity**, needing only first derivatives. This approximation is well-justified for residual-driven costs – it effectively linearizes the model, solving the quadratic subproblem at each iteration. As long as this linearization is valid, Gauss–Newton converges almost as fast as Newton’s method but with more stability. If the linear model is poor (residuals large or highly nonlinear), one must take precautions (damping, small step sizes, better initial guess). In practice, the combination of Gauss–Newton with damping (Levenberg–Marquardt) is a workhorse in bundle adjustment and ICP because it balances speed and robustness.

## 7. Role of the Hessian’s Condition Number in Optimization

The **condition number** of the Hessian  $H$  – roughly the ratio between its largest and smallest singular value (or eigenvalue, if symmetric) – plays a critical role in the performance of optimization algorithms and in the interpretation of uncertainty:

- **Convergence speed and directionality:** A well-conditioned Hessian (low condition number, eigenvalues of similar scale) means the quadratic approximation of the cost is nearly isotropic. Algorithms like Gauss–Newton or Newton’s method will then take steps that approach the optimum from any direction at similar rates. If  $H$  is ill-conditioned (high condition number), the error contours are elongated like a valley. The optimizer may experience the classic zig-zag: quickly correcting along the steep directions but slowly inching along the shallow direction. In extreme cases, the shallow direction might cause the linear solver to produce a large step that overshoots (if not damped) or a very tiny step (if damped), either of which slows convergence. For example, if one eigenvalue is  $1000\times$  larger than another, the update calculation  $\delta\omega = H^{-1}g$  will have components that differ by that factor, making it hard to choose a single step size that suits both. Gradient descent in such a scenario is slow (needs very small learning rate to handle the steep direction), whereas Gauss–Newton is better but still might need damping.
- **Step stability:** Solving the normal equations  $H\delta\omega = -g$  (with  $g = -J^T r$  the gradient) when  $H$  is ill-conditioned can lead to numerical issues. A very high condition number means the matrix is nearly singular; finite numerical precision can amplify errors. This might result in an inaccurate  $\delta\omega$  solution where the component along the small eigen-direction is noisy or undefined. In SLAM and point cloud alignment, one symptom is jitter or instability in the solution for that axis. To improve stability, one may use more robust linear solvers (like SVD or QR factorization) which can handle rank-deficient cases by truncation. Indeed, one approach to **degeneracy detection** is to monitor the minimum eigenvalue of  $H$  or its condition number – if it falls below a threshold, the system is declared ill-conditioned and the solver might regularize or hold back certain updates.
- **Regularization and damping:** If the condition number is large, adding a damping term  $\lambda I$  (Levenberg–Marquardt style) effectively increases the smallest eigenvalues more (relative to the largest), thus improving the condition number. This makes  $H + \lambda I$  easier to invert and yields more conservative updates. The LM algorithm adapts  $\lambda$  such that when the system is ill-conditioned or far from linear, the step is small (close to gradient descent which doesn’t trust the Hessian as much). As the system becomes more conditioned (near optimum),  $\lambda$  is decreased, letting Gauss–Newton take larger steps. Another approach is **Tikhonov regularization**, which is essentially the same idea: solve  $(H + \lambda I)\delta\omega = -g$  to avoid uncontrolled

swings in poorly observed directions. This also has an interpretation in Bayesian terms (adding a prior on  $\omega$  to cap uncertainty).

- **Uncertainty quantification:** As noted earlier,  $H^{-1}$  (or scaled version) gives the covariance of the estimate. A high condition number means this covariance matrix is very elongated – huge uncertainty in some direction. For diagnosis, one can look at the eigenvector corresponding to the smallest eigenvalue of  $H$ : this is the axis about which the solution is least certain. In a robot localization context, for instance, if  $H$  is ill-conditioned, it might indicate the robot’s orientation about a certain axis is not observable (e.g., all points lie on a plane, so rotation around the normal is hard to determine). Recognizing this, the algorithm designer might add an additional sensor or prior to constrain that direction, or at least be aware that the estimate angle is unreliable. Some SLAM systems avoid adding constraints (factors) that lead to nearly singular Hessians (they detect degeneracy to maintain stability).
- **Diagnostics and improvements:** Monitoring the condition number of  $H$  per iteration can alert one to a problematic configuration. If it’s above a certain threshold, the solver might:
  - Increase damping (LM does this automatically).
  - Re-initialize or try a different initial guess (in case we are stuck in a weird configuration).
  - In mapping problems, **add constraints**: e.g., if pure ICP on a planar wall is degenerate for rotation about the normal, combining it with a small IMU rotation measurement can improve conditioning.
  - Use robust methods like **Truncated SVD**: solve for  $\delta\omega$  by truncating tiny singular values of  $H$  (essentially ignoring the unobservable component). This yields a solution in the subspace that is well-constrained.
  - In point selection (if we can choose points as in an active sensing scenario), pick points that improve conditioning as discussed in Section 2. In fact, degeneracy-aware algorithms sometimes resample point correspondences until the condition number of the resulting Hessian is below a threshold, explicitly enforcing good conditioning.

**Concrete example (condition number effect):** Suppose we align a point cloud of a long, thin rod. If we only have that rod’s points, rotation about the rod’s axis is nearly unobservable – Hessian eigenvalues might be  $[100, 100, 0.01]$  for rotations about two perpendicular axes vs the rod axis. The Gauss–Newton step for rotation about the rod axis will be very hesitant or noisy. If we add another point off the rod’s axis, the smallest eigenvalue might rise to, say, 5. Now condition number improves from  $100/0.01 = 10000$  to  $100/5 = 20$ , a huge improvement. The optimizer will converge much more surely. This underscores the importance of conditioning: it’s tied to both algorithmic performance and the fundamental physics of observability.

In summary, the condition number of the Hessian is a critical indicator of

**problem difficulty.** Low condition numbers mean fast, stable convergence and low estimation uncertainty. High condition numbers mean slow convergence, potential instability, and ambiguous results. Practically:

- Always check if your Hessian (or information matrix) is ill-conditioned or singular; it might indicate degenerate geometry or insufficient data.
- Mitigate poor conditioning by adding information (more diverse points or a prior), or by using algorithmic safeguards (damping, regularization, specialized solvers).
- Use the Hessian’s eigenvalues as a diagnostic tool for which degrees of freedom are weakly constrained – this can guide sensor placement or further iterations in a larger system.

**Conclusion:** The Gauss–Newton Hessian  $H = J^T J$  for 3D rotation encapsulates the geometry of the point configuration, the certainty of the estimate, and the difficulty of the optimization. A derivation gives  $H = \sum_i (\|x_i\|^2 I - x_i x_i^T)$ , linking to the inertia matrix of points. This form immediately tells us how to choose points for maximal observability (no degeneracies) and how to recognize when rotations are unobservable (rank-deficient  $H$ ). The Hessian’s properties (PD, eigen spectrum, condition number) connect to optimality conditions and to the Fisher information, reinforcing that a well-conditioned Hessian implies a precise, reliable solution. When implementing these optimizations, using left perturbations keeps the math straightforward and consistent, and employing Gauss–Newton (with damping if needed) strikes a good balance between complexity and performance, relying on small residuals to justify neglecting second-order terms. Finally, always mind the condition number of  $H$ : it affects everything from convergence speed to the physical interpretability of your solution’s uncertainty. By monitoring and managing it, one can diagnose issues and ensure robust 3D pose optimization in practice.

**Sources:** The above analysis synthesizes standard results from optimization and robotics literature. The skew-symmetric derivative identity and Hessian form are textbook in hand–eye calibration and point cloud registration. The Fisher information connection to Gauss–Newton is well-known in estimation theory. Degeneracy and conditioning issues in ICP and SLAM have been widely discussed, e.g. degeneracy detection via Hessian eigenvalues and improving conditioning by point selection. These principles form the backbone of practical algorithms for pose estimation and SLAM.